

FOR PRODUCT MANAGERS WHO SHIP WITH AGENTS

Supaprod

agents that know what to build,
ship it, and guide the next call.

PRE-SEED BRIEFING · AUGUST 2026 · CONFIDENTIAL

· Three minutes · press →

One system, three layers.

01 · THE DIRECTOR

It tells you what to build.

Reads your feedback, data, competitors, and market. Tells you what is worth building next, evidence attached. Your product taste becomes a system.

02 · THE OPERATING SYSTEM

It runs the whole lifecycle.

Discover to learn, one governed loop. Agents do the work in your own stack. You make the calls.

03 · THE BRAIN

It learns, and it guides.

Every decision recorded with its evidence, then graded against what happened. It compounds: next time it tells you what is right, and warns you before you repeat what was wrong.



Every craft got its AI-native home.

Product decisions still live in Slack.

CODE	Claude Code · Cursor · Codex · v0 · Lovable	AGENT-RUN
DESIGN	Figma · Framer	AGENT-RUN
DOCS	Notion AI · Gemini	AGENT-RUN
WORK	Linear · Jira	SYSTEM OF RECORD
DECISIONS	heads, threads, scrollbar	NO HOME

“So why did we decide on X? Cue hours of finding that Slack conversation from months ago.”

Top-voted thread · r/ProductManagement · 480 points

Nobody knows what to build, so teams build on gut: **80% of shipped features are rarely or never used** (Pendo). **The one record a team needs most is the one nobody keeps.**

AI made building cheap. Deciding what to build is the bottleneck.

WHAT JUST HAPPENED ONE LAYER BELOW US

\$2B+

Cursor annual revenue, four years after launch

CNBC · FEB 2026

75%

of new code at Google is now written by AI

SUNDAR PICHAI · APR 2026

8 mo

for Lovable to reach \$100M ARR. Now past \$500M

TECHCRUNCH · JUN 2026

Every engineer got an AI pair. **The person deciding what they build got a chatbot.**

Judgment has no compiler.

Code has a fast oracle: tests pass or they fail. That tight loop is why AI coding exploded and commoditized. A product decision waits months for its verdict. **Judgment never got its loop.**

signal → decision → evidence → shipped → outcome → was the reasoning right?

The track record is the compiler for judgment.

Twelve steps. Nineteen minutes.

You made two calls.

A morning on **Supaproduct**: it spots the problem, proposes the fix, you approve, agents ship, and D+14 grades the call.

08:02	IT NOTICES	your two largest accounts hit the same wall this week, and renewal season starts in six	THE DIRECTOR
09:14	IT PROPOSES	fix the export flow first: here is the revenue at stake, and the two cheaper options it rejected	THE DIRECTOR
09:14	IT LEARNS	the last time a big account hit friction and you waited, they churned. Here is the evidence	THE BRAIN
09:16	YOU DECIDE	one click. The only human moment in the run	
09:17	AGENTS SHIP	design, code, tests. Live behind a safety flag the same day	THE OPERATING SYSTEM
09:17	IF IT FAILS	any failed check stops the run and returns it to your gate. Nothing ships without you	THE OPERATING SYSTEM
D+14	IT GRADES	exports up, tickets gone, both accounts renewed. Verdict saved, so the next call starts sharper	THE BRAIN

LIVE RUN · AGENT MESH		3 / 5 WORKING
● Scout	clustered 3 signals	DONE
● Architect	spec locked, 4 criteria	DONE
● Designer	screens drafted from spec	RUNNING
● Builder	code written, checks green	QUEUED
● Sentry	watching D+14 outcome	QUEUED

Agents ran twelve steps in nineteen minutes. You made two calls. **Every one is on the record.**

It learns your taste, not just your tasks.

Verdicts take weeks, so the record accrues in calendar time. **What no one can reconstruct afterwards is what you believed before.** A forecast is not an artifact. It exists only if something captured it at the moment you decided, so a competitor starting next year starts at zero, next year.

Quarter two starts smarter than quarter one.

Two budgets are converging on one layer.

THE TOOLS BUDGET · PM SOFTWARE

\$3.8B today

Steady, unglamorous, growing 14% a year, and about to be re-platformed by agents.

MARKWIDE RESEARCH · MAY 2026

THE AGENT BUDGET · UNGOVERNED

5% → 40%

Share of enterprise apps shipping task-specific AI agents, 2025 to end of 2026. Every one creates decisions nobody scores, and unscored decisions never improve.

GARTNER · AUG 2025

2.6M people hold the product manager title on LinkedIn. The wedge customer: the founding PM already shipping with agents. Land at \$20 a month, expand to the team, and **grow on decision volume, not headcount.**

LAUNCH · SEP 2026

Public and self-serve from day one. Every consequential call human-gated.

SCALE · IMMEDIATELY

No slow-burn roadmap. Individual PMs to teams to orgs, meter on, aggressive from the first week.


ENTERPRISE · FAST FOLLOW

Governance and audit. The track record becomes the org's institutional memory.

AUTONOMY EXPANDS PER LANE, EARNED ON TRACK RECORD. STRATEGIC CALLS STAY HUMAN, PERMANENTLY.


The tools budget is billions.
The work budget is hundreds
of billions.

TAM ~~\$300B+~~ / yr THE PM WORK BUDGET



What 2.6M PMs are paid to do: know what to build, ship it, learn from it. **Agent swarms move that spend from headcount to systems.**

SAM \$2B → \$12B / yr LAUNCH PRICING → VALUE PRICING



MOTION 1 · TRANSFORM 650K existing product teams MOTION 2 · CREATE 500K new agent-native orgs by 2030

SOM ~\$47M ARR THE ENTRY WEDGE · TODAY



The agent-native tenth of existing teams, at launch pricing. **One layer below, Cursor crossed \$2B ARR by year four on this same curve.**

Priced on outcomes.

Agents do not have seats.

A credit pays for a finished result: a closed decision loop. Never tokens, never chat, never seats. A teardown runs 15 to 40 credits; idea to shipped, 150 to 400. **The price of an afternoon of coordination.**

1 · LAND

Free teardown of your riskiest idea. No card.

2 · CONVERT

Pro \$20/mo solo · Business \$50/mo team, credits included.

3 · EXPAND

Heavy shipping months buy top-ups. Spend grows toward a tenth of one PM's cost per team, tracking decisions closed, not headcount.


4 · ENTERPRISE

SSO, compliance, audit, BYOK. Governance is what enterprises pay for; the track record is the audit trail.

CREDITS POOL AT THE ACCOUNT LEVEL, NEVER PER SEAT, SO SHARED MEMORY IS NEVER TAXED

Seat pricing shrinks as PM work automates. **Credit pricing grows with every decision the org makes.**

Everyone owns a step. Nobody owns the loop.



THE VERTICAL AXIS COMPOUNDS: EVERY GRADED OUTCOME MOVES US UP. THAT IS THE MOAT

None of them records whether the call was right. That needs the whole loop under one roof. It cannot be bolted on, and the belief you held going in cannot be recovered after the fact.

Our real competitor is the **stitched stack**: five tools and the PM as the glue.

Not a wrapper on a frontier model. The models are interchangeable parts; **the system is ours: the loop, the gates, the track record**. Our engine runs the best model for each job, in your

Built. Running. Opening.

- Listened first: pain research across PM communities, practitioners, and the stitched-stack workflow DONE
- Prototyped, built, rebuilt. Kept what survived contact with real use DONE
- The full loop, discover to learn RUNNING TODAY
- Self-serve signup LIVE
- Fine-tuning the experience, folding in early feedback IN PROGRESS
- Early access invites, from a named list NEXT
- Public launch SEPTEMBER 2026

Pre-revenue, pre-launch, said plainly. **The working engine is the proof; the cohort is the next one.**

Built by user zero.



Rohit Gajaraj

FOUNDER · A DECADE IN PRODUCT · FULL-TIME ON SUPAPROD

✉ founder@supaproduct.ai

in LinkedIn

EXPERIENCE

Intellect	Senior AI product manager at Intellect, a leading BFSI technology OEM serving 200+ financial institutions across 70+ countries. BANGALORE, INDIA
Infineon	Product manager at a leading global semiconductor company. MUNICH, GERMANY
ISRO	Associate product manager, communication systems at India's national space agency. Systems that flew, built at 21, where a mistake is unrecoverable. BANGALORE, INDIA

EDUCATION

TUM	MBA, TUM School of Management, one of Europe's top business schools, at the Technical University of Munich. MUNICH, GERMANY
-----	--

WHY THIS FOUNDER

Ten years being the glue this product replaces: carrying context across a dozen tools and answering “why did we decide this” from memory, months later. Shipped where a wrong call is unrecoverable: space, silicon, banking. **He built the thing he needed, and runs on it every day.**

BUILT FOR AGENTS TO RUN

Supaproduct is built the way it expects product teams to work: **one human on judgment, an agent swarm on execution**, every step gated. Solo, in weeks. The proof of the method is the product itself.

Agents do the work.
You answer for it.

Supaprod is how you answer.

- TELLS YOU WHAT TO BUILD
- RUNS THE WORK
- LEARNS, AND GUIDES THE NEXT CALL

Every product org is about to run on agents. The winners will know **what to build**, ship it without drag, and **get sharper with every call**. That is not a feature on a tracker. It is a new operating system for product. **We are building it.**

Sign up for the beta → supaprod.ai

Investor relations · investors@supaprod.ai · launching September 2026



Who is who, and what they do not keep.

RELATION	WHO	WHAT THEY DO	WHAT THEY DO NOT
RACE	Productboard Spark (agentic relaunch Jun 2026) · Atlassian Product Collection (May 2026) · Samepage Signals (\$4.85M seed, Jun 2026) · Brief · Notion AI · Linear's AI agenda	AI on top of feedback, roadmaps, tickets, and docs. Sense and summarize; some draft and prioritize. Each owns a part of what we do.	None ships the work or grades the outcome. The loop never closes, so nothing learns. No decision-and-outcome record.
ABSORB	ChatPRD (50K+ users) and AI spec writers	Excellent single artifacts: PRDs, specs. Proof that PMs pay for AI help.	A spec is one step of one stage. No loop, no record, no verdicts.
ENGINES	Frontier models via API (Claude, GPT, Gemini) · era tools: Claude Code , Cursor , Lovable , Devin , Windsurf , JetBrains agents	Raw building capability. Our own build engine runs the same models inside the governed loop, in your repo.	Capability is not a system. No gates, no track record, no accountability for what was built or why.
SENSE SILOS	Dovetail · Enterpret · live channels: Intercom · Zendesk · Slack	Feedback and research repositories, and the channels customers actually speak in. Strong at remembering what users said.	Insight stops at the document. Nothing ships, nothing is graded, nothing learns.
RAILS	GitHub · Linear · Jira · Notion · Figma · CRM: Salesforce · HubSpot	Systems of record for code, tasks, docs, design, and who the customer is. We read and write them.	No home for decisions, evidence, or outcomes. The why still dies in threads and pages.
MEASURE	Amplitude · Mixpanel · PostHog · Statsig · LaunchDarkly	Where the result of a shipped bet actually lands. We read them; this is the raw material of a verdict.	Hold the number, never the decision that caused it. Nobody joins the outcome back to the call, so nothing re-ranks the next bet. This is the gap we are.

We touch six surfaces and compete on one: whether a shipped bet's outcome is verified and fed back into the next decision. Everywhere else we read, absorb, or run on top of. Full competitive Q&A brief available on request.

Every figure in this deck, sourced.

\$1B → \$2B+	Cursor ARR, Nov 2025 to Feb 2026; \$29.3B valuation	CNBC
\$100M / 8 mo	Lovable ARR ramp; \$500M+ by Jun 2026	TECHCRUNCH
75%	of new Google code AI-generated, up from 50% six months earlier	THE VERGE · APR 2026
20M · 90%	GitHub Copilot all-time users; share of Fortune 100 using it	MICROSOFT FY25 Q4
80%	of shipped features rarely or never used, across 615 products	PENDO FEATURE ADOPTION REPORT
5% → 40%	enterprise apps with task-specific AI agents, 2025 to end of 2026	GARTNER · AUG 2025
\$3.8B → \$12.6B	PM software market, 2026 to 2035, 14.2% CAGR	MARKWIDE RESEARCH · MAY 2026
2.6M+	people holding the product manager title on LinkedIn	PRODUCTIFY · MAY 2026
\$300B+ / yr	TAM math: 2.6M PMs (LinkedIn) × ~\$115K average fully loaded cost (salary, benefits, overhead, blended globally) = \$299B of product-management work paid as headcount today	INTERNAL · TRANSPARENT MATH
\$2B → \$12B	SAM math: 650K existing teams (2.6M PMs ÷ 4) + 500K new agent-native orgs by 2030 = 1.15M orgs. At launch pricing (\$150/mo): ~\$2B a year. As autonomy earns value pricing (~\$10K a year, a tenth of one PM's cost): ~\$12B a year, about 4% of TAM	INTERNAL · TRANSPARENT MATH
~\$47M ARR	SOM math: 10% of existing teams are agent-native today (Gartner trajectory: 5% to 40%) = 65K teams × ~\$60 a month blended launch pricing = \$47M a year, before value pricing	INTERNAL · LABELED SCENARIO